



THE UNIVERSITY OF  
**CHICAGO**

Office of Research and  
National Laboratories  
Research Computing Center

# Running Jobs on RCC Systems Using the SLURM Scheduler

Teodora Szasz, PhD

Computational Scientist

October 15th, 2020

# Agenda

- The RCC Midway compute systems
- Using SLURM (Simple Linux Utility for Resource Management) to submit jobs to the RCC Midway systems

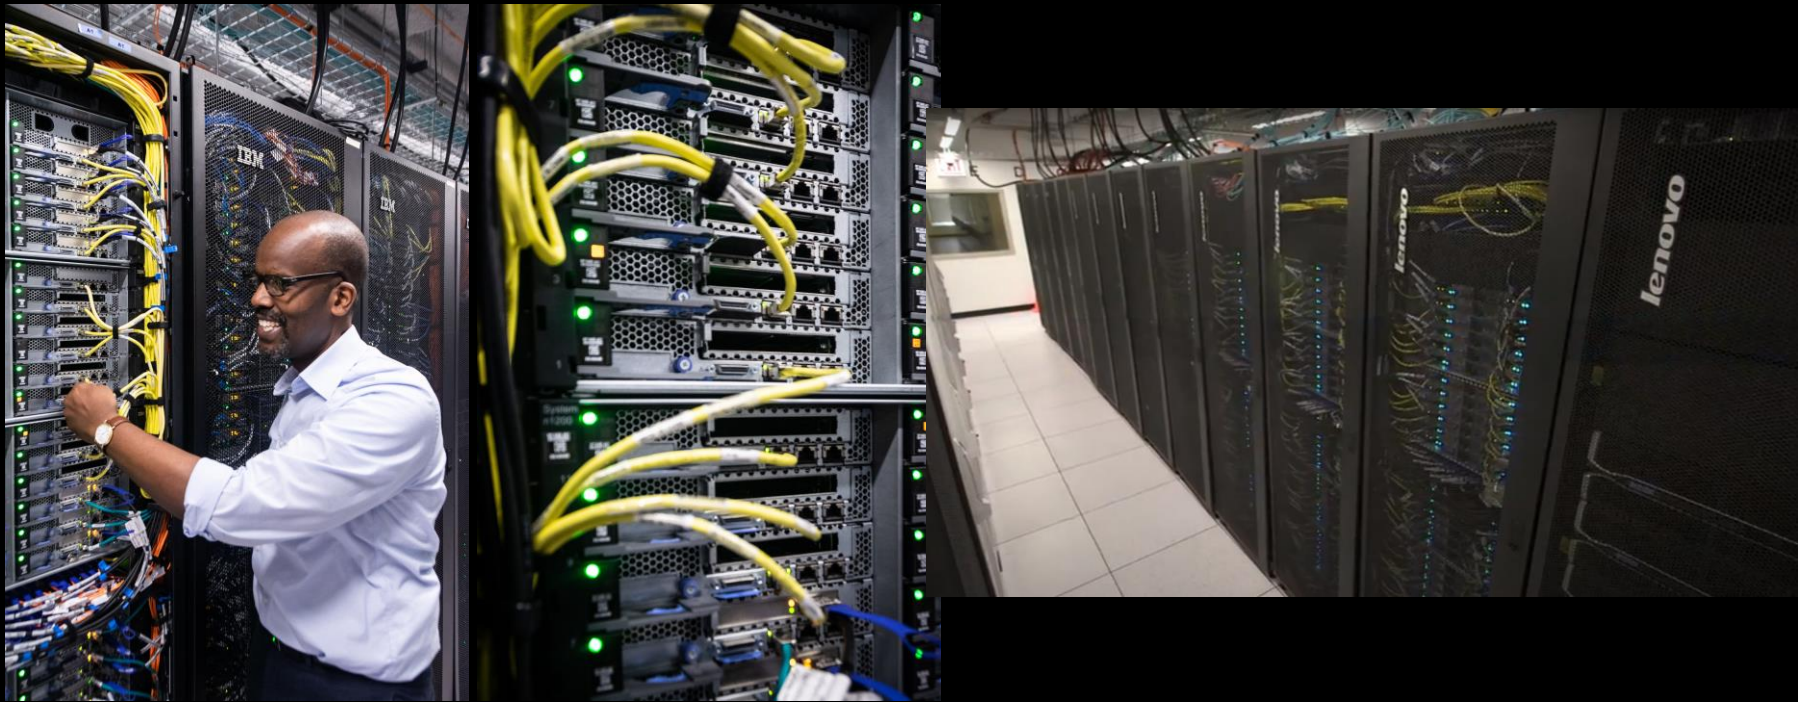
## => Goals of this workshop

- Learn the basics of SLURM
- Submit jobs on Midway HPC systems
- Understand priority queue
- Check if your jobs succeeded/failed
- Know how to fix common errors



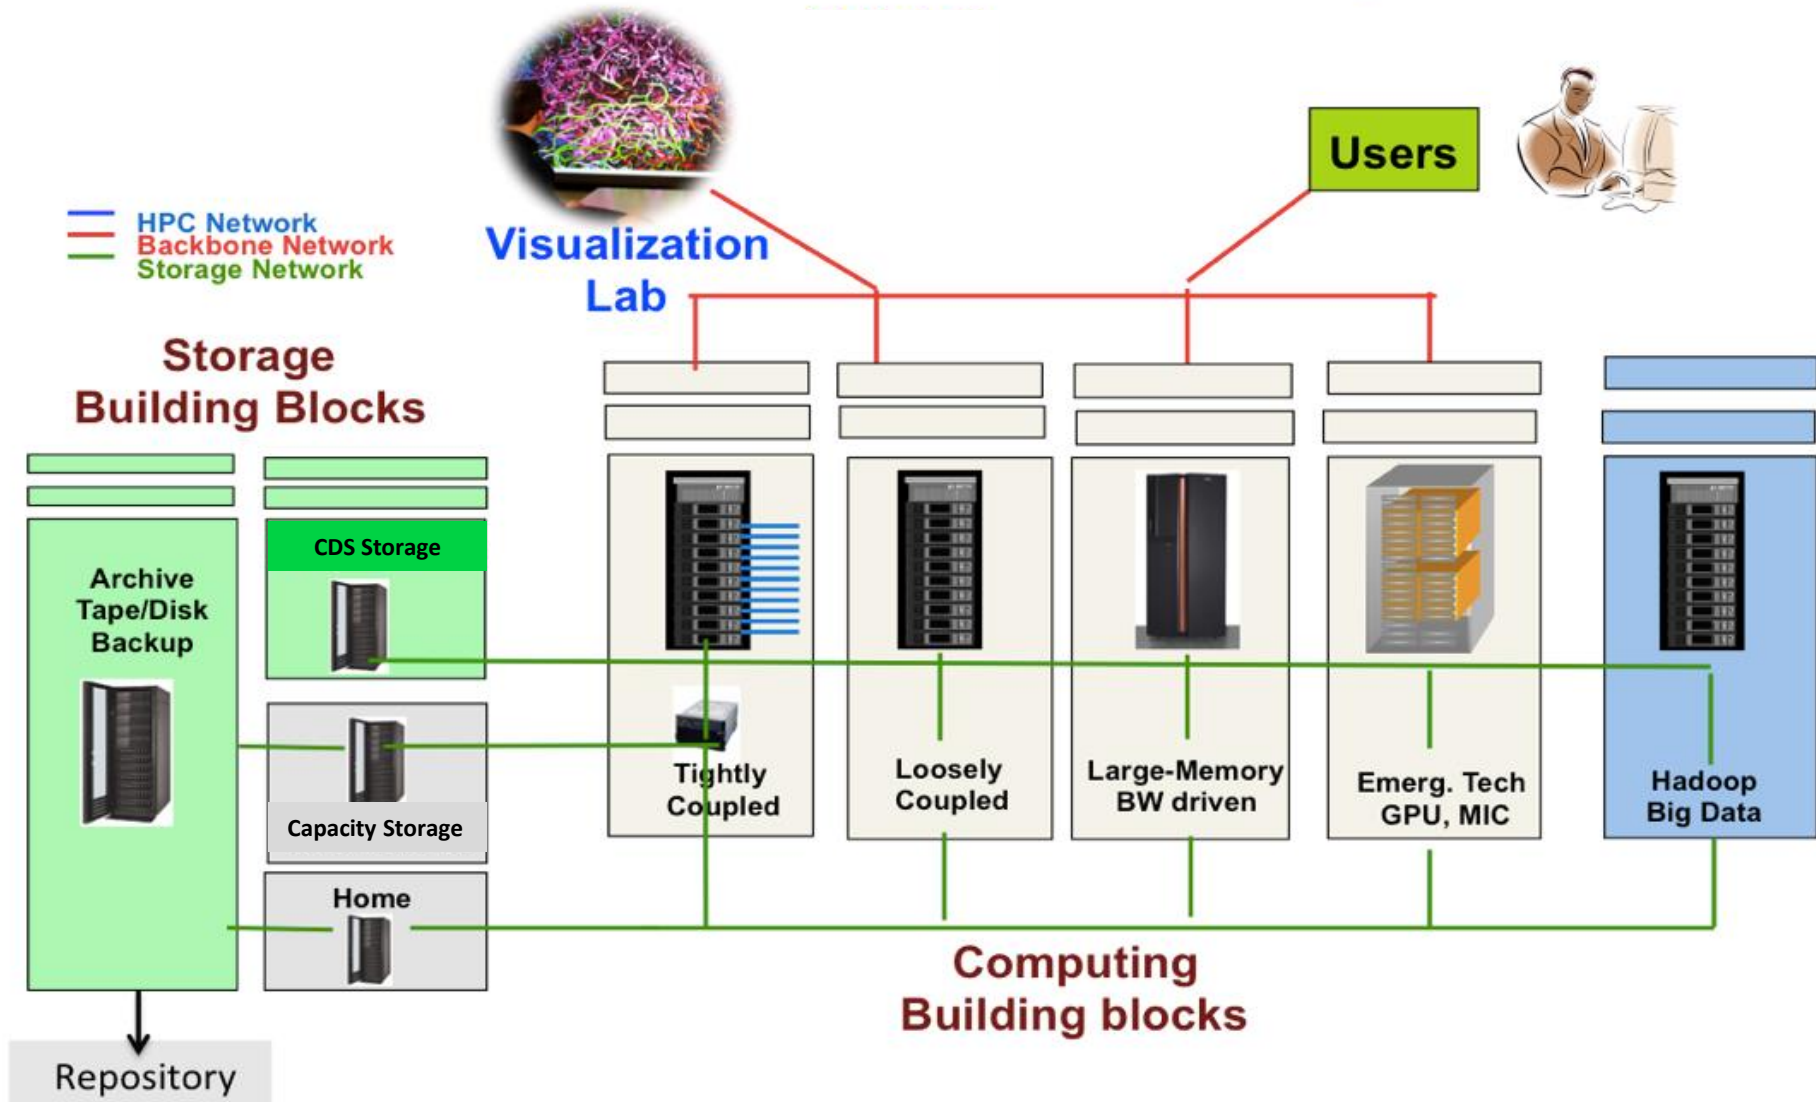
# Understanding the RCC Compute Ecosystem

Midway is a constellation of many compute systems and storage with various architectures coupled together in one system.



SLURM is the software used to manage the workload on Midway

# Hardware Infrastructure building blocks



# Some definitions

- A **processor** is a small chip that responds to and processes the basic instructions that drive a computer. The term *processor* is used interchangeably with the term **central processing unit (CPU)**
- **Core:** The smallest compute unit that can run a program
- **Socket:** A compute unit, packaged as one and usually made of a single chip often called processor. Modern sockets carry many cores (10, 14, or 20, 24, 28, etc. on most servers)
- **Node:** A stand-alone computer system that contains one or more sockets, memory, storage, etc. connected to other nodes via a fast network interconnect.

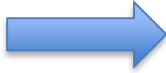




# From a CPU to a Cluster



**CPU**  
(Has Multiple cores)



**Socket**  
(Can have Multiple CPUs)



**Motherboard**



**Compute Node**

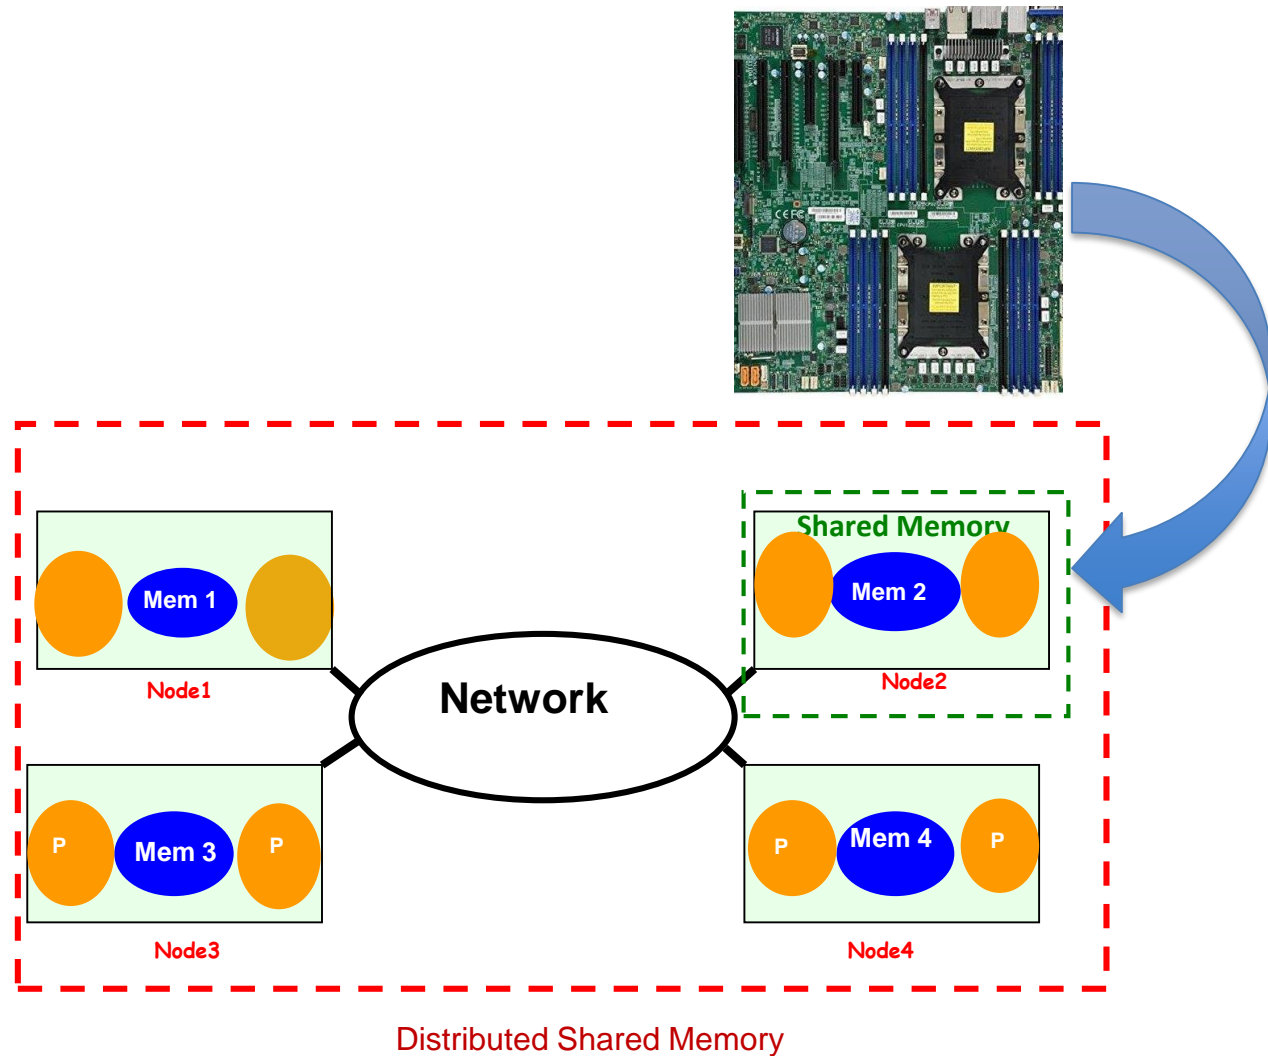


**Rack**



**Cluster**

# Distributed vs Shared Memory Systems



# The RCC compute cluster

## Computing hardware

### **367 nodes total**

- **342** tightly coupled Broadwell nodes ( 10,360 cores)  
Two intel E5-2680v4 processors per node (14 core/proc)  
155 nodes have EDR network card  
187 nodes have FDR network card
- **6** NVidia Tesla K80 GPU nodes (4 GPU cards/node)
- **5** large shared memory nodes (512GB each)
- **14** dual socket loosely-coupled Broadwell nodes

## **Cluster Partnership Program: 1000+ nodes**

- **900+** tightly coupled infiniband nodes
- **20+** Big memory nodes
- **100+** Nvidia GPU nodes



**RCC Manages 1515 nodes  
(35,904 cores)**





# RCC Hardware: Midway2 (Shared)

Cores(Nodes)	CPU	Memory	Network
5236 (187)	2.4GHz Intel Broadwell	64GB	EDR 100Gbit/s InfiniBand
4,340 (155)	2.4GHz Intel Broadwell	64GB	FDR14 56Gbit/s InfiniBand
140 (5)	2.4GHz Intel Broadwell	512GB	FDR14 56Gbit/s InfiniBand
504 (14)	2.4GHz Intel Broadwell	64GB	1 GigE

Cores (Nodes)	CPU Accelerator	Memory	Network
168 (6)	2.4GHz Intel Broadwell NVIDIA Tesla K80	128GB	FDR14 56bit/s InfiniBand



# RCC Hardware: Midway2 (Shared)

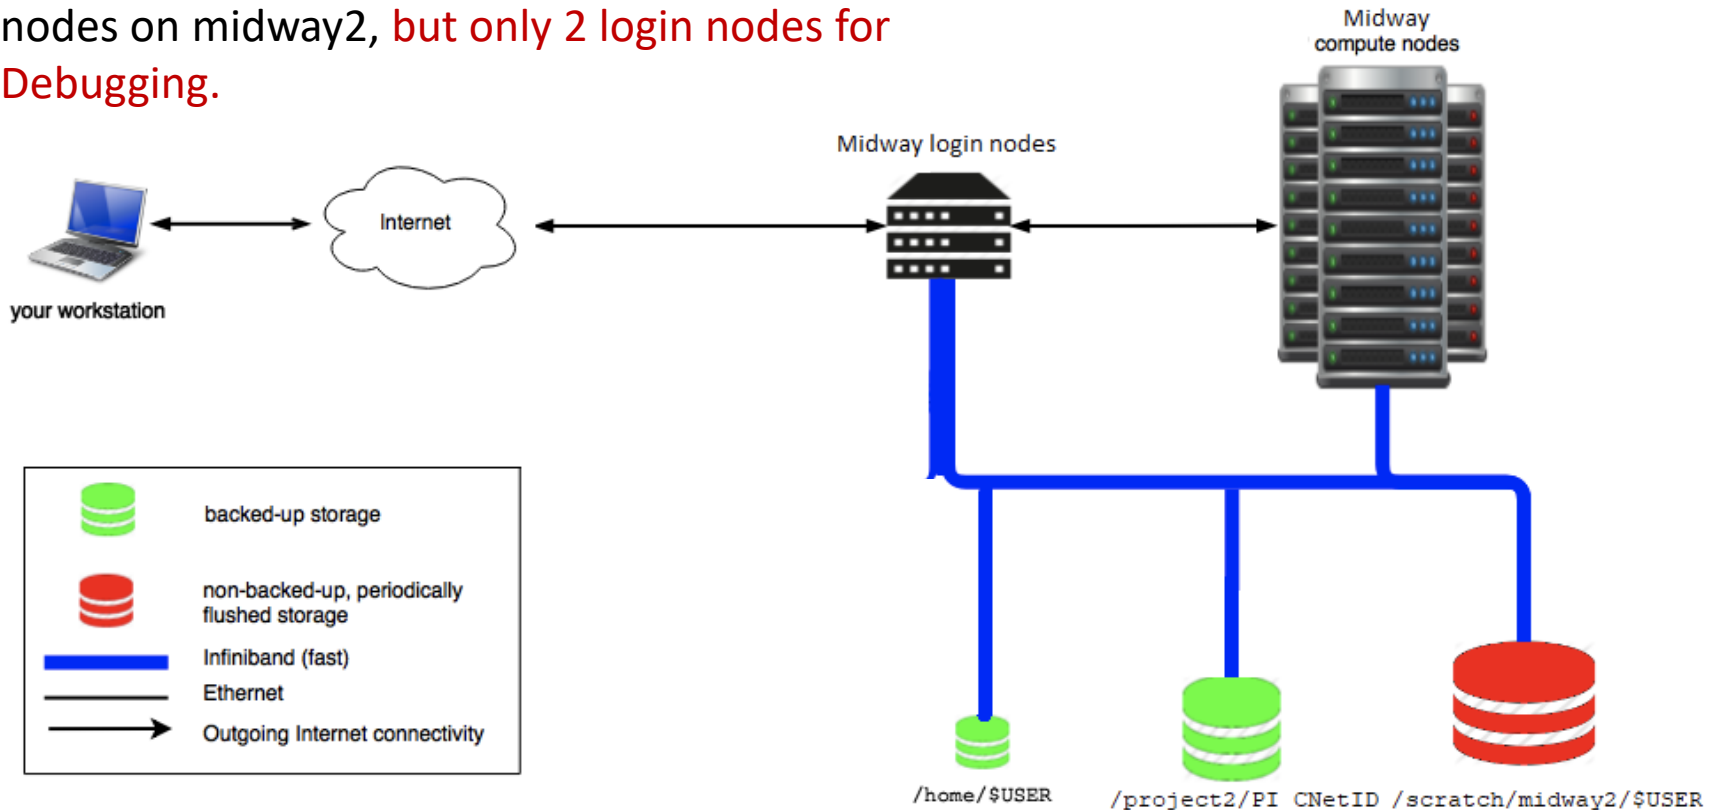
Cores(Nodes)	CPU	Memory	Slurm Queue
5180 (185)	2.4GHz Intel Broadwell	64GB	broadwl
5180 (185)	2.4GHz Intel Broadwell	64GB	broadwl
168 (6)	2.4GHz Intel Broadwell	512GB	bigmem2
504 (18)	2.4GHz Intel Broadwell	64GB	broadwl-lc

Cores (Nodes)	CPU Accelerator	Memory	Queue
168 (6)	2.4GHz Intel Broadwell NVIDIA Tesla K80	128GB	gpu2



# Schematic of the Midway Cluster

There are about more than 1300+ nodes compute nodes on midway2, **but only 2 login nodes for Debugging.**



# Midway2 Storage

## High Capacity storage: /project2

- **3.8 PB** of storage
- Backed up to tape system
- **7 daily and 4 weekly snapshots** located at /snapshots/project2
- 7 day grace period on over quota

## High Performance storage: /scratch/midway2

- **190 TB** usable
- Not backed up
- 100 GB user soft quota
- 30 day grace period on over quota



## Home directory space: /home

- **61 TB** of capacity
- Each user has 30 GB quota
- 7 day grace period on over quota
- **7 daily and 2 weekly snapshots** located at /snapshots/home



# How to Run jobs at RCC

- Interactively
- Submitting jobs to a queue using SLURM



# Repository

Please clone the Github repo at:

[https://github.com/rcc-uchicago/SLURM\\_WORKSHOP.git](https://github.com/rcc-uchicago/SLURM_WORKSHOP.git)

```
[tszasz@midway2-login1 Intro_SLURM]$ git clone https://github.com/rcc-uchicago/SLURM_WORKSHOP.git
```

# A Key point to remember

There are about more than 1300+ nodes compute nodes on midway2, but only 2 login nodes.

This means you are sharing the login nodes with many other users at once. Running intensive programs on the login nodes causes the login nodes to be slow for all other users.

- login nodes are for editing files, compiling, moving files, changing permissions, and other non-intensive tasks.
- We recommend to use *sinteractive* for interactive runs
- For long running jobs => submit them to the priority queue through SLURM

# Running Interactive jobs

- Login directly to a node
  - Login to `midway2.rcc.uchicago.edu`
- Run interactively using `sinteractive`  
`sinteractive --time=01:00:00 --nodes=1`  
`--ntasks=2 --mem-per-cpu=1000 --`  
`partition=broadwl --reservation=slurm-`  
`workshop`



# Summary of partitions and nodes on Midway2

**sinfo -s**

```
[rajshukla@midway2-login1 ~]$ sinfo -s
PARTITION    AVAIL  TIMELIMIT  NODES(A/I/O/T)  NODELIST
cron          up    infinite    0/4/0/4          dali-login[1-2],midway-login[1-2]
westmere      up    infinite    0/28/1/29        midway[002-030]
sandyb        up    infinite    49/24/2/75       midway[044,069-073,089,109,112,193-197,216-226,398-414,417-445,448-451]
mfj           up    infinite    1/11/0/12        midway2-[0489-0500]
test          up    infinite    0/1/0/1          midway397
cobey         up    infinite    16/0/1/17        midway2-[0217-0220,0409-0410,0427-0436],midway2-bigmemo5
jnovembre     up    infinite    4/4/0/8          midway2-[0401-0408]
tas1          up    infinite    0/19/5/24        midway[783-806]
gpu           up    infinite    0/8/1/9          midway[230-232,493-494],midway-134-[01-04]
viz           up    infinite    0/1/0/1          midway229
mic           up    infinite    0/0/2/2          midway-mic[01-02]
sepalmer      up    infinite    0/3/0/3          midway[453-454,590]
kicp          up    infinite    0/20/11/31       midway[159-188,191]
kicp-long     up    infinite    0/20/11/31       midway[159-188,191]
kicp-ht       up    infinite    0/4/0/4          midway[151-152,189-190]
surph         up    infinite    0/7/7/14         midway[143-150,153-158]
surph-large   up    infinite    0/27/18/45       midway[143-150,153-188,191]
xenon1t       up    infinite    14/2/0/16        midway2-[0411-0426]
```

⇒ NODES(A/I/O/T) : Nodes (Allocated/Idle/Other/Total)

e.g: To check the broadwl partition: **sinfo -s | grep broadwl**

```
[tszasz@midway2-0073 Intro_SLURM]$ sinfo -s | grep broadwl
broadwl*      up    infinite    214/133/13/360   midway2-[0002-0133,0137-0190,0221-0280,0282-0299,0301,0311-0398,0461-0466,0612]
broadwl-lc    up    infinite    12/0/2/14        midway2-[0203-0216]
```

# SLURM: Some key terms to remember

A **job defines** the resources you are using and the code you are running

The **queue** in SLURM is all RUNNING and all PENDING jobs  
To see every job in the queue, use the command:

```
squeue
```

To see your jobs in the queue

```
squeue -u <cnetid>
```

or

```
myq
```



# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash

# Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

This **#!** is a shebang

It tells operating system to use  
**/bin/bash**  
with this script

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

# is a comment

everything after # is ignored by  
bash

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

**#SBATCH** is a directive

It is a comment in Bash

**#SBATCH** is only relevant to  
SLURM:

```
sbatch my_script.sh
```

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash

# Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

**#SBATCH** is a directive

It is a comment in Bash

**#SBATCH** is only relevant to slurm:  
sbatch my\_script.sh

To comment out directives, break the pattern, e.g.

##SBATCH

# SBATCH



# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash

# Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

Instructions for SLURM must go at the top of the script

Any #SBATCH lines you put after your program will be ignored

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash

# Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

Slurm has some variables you can use. %j is the job number. When the job runs %j will be expanded to the job number. In this example %j is used in the output file and error file names:

MyJob-13571056.out  
MyJob-13571056.err

%j is unique. By using %j in your filenames you guarantee a unique file name, which means you won't accidentally overwrite previous output.

# What goes into a batch script?

A batch script is list of instructions for SLURM.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

=> Time your job is allowed to run

```
#SBATCH --nodes=1
```

=> Number of nodes to run on

```
#SBATCH --ntasks-per-node=1
```

=> Number of cores on each node to use

```
#SBATCH --mem-per-cpu=2000
```

=> Memory per cpu => 2000Mb or 2Gb

```
#SBATCH --job-name=MyJob
```

=> Name of the job.

```
#SBATCH --output= MyJob-%j.out
```

=> Job output file behaves as stdout for the code.

```
#SBATCH --error=MyJob-%j.err
```

=> Error file. behaves as stderr for the code.

```
module load <module name>
```

=> Load any modules you need for your application

```
#Run your code
```

=> run the code you want

# Running batch jobs using a Submission Script

- A simple job submission script (saved as python.sbatch):

```
SLURM { #!/bin/bash
        #SBATCH --job-name=first_python_job
        #SBATCH --output=first_python_job_%j.out
        #SBATCH --error=first_python_job_%j.err
        #SBATCH --nodes=1
        #SBATCH --ntasks-per-node=1
        #SBATCH --mem-per-cpu=2000M
        #SBATCH --partition=broadwl
        #SBATCH --reservation=kicpworkshop-cpu
Your Job { #SBATCH --time=00:30:00
            module load python
            python hello_world.py
            echo "job finished at `date`"
```

- To submit the above script:
  - **sbatch** python.sbatch



Please do the Exercise;  
Ex-1 in Repo



# How to submit OpenMP jobs?

```
#!/bin/bash

#Here is a comment
#SBATCH --time=1:00:00
#SBATCH --partition=broadwl
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=8
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
export OMP_NUM_THREADS=8
#Run your code
./my_executable
```

Specify number of cores > 1.

OMP\_NUM\_THREADS is an environment variable.

Please do the Exercise;  
Ex-2 in Repo

# How to submit Parallel MPI jobs?

```
#!/bin/bash
```

```
#Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
#SBATCH --partition=broadwl
```

```
#SBATCH --nodes=4
```

```
#SBATCH --ntasks-per-node=8
```

```
#SBATCH --mem-per-cpu=2000
```

```
module load openmpi
```

```
module load <module name>
```

```
#Run your code
```

```
mpirun ./my_executable
```

Specify number of nodes > 1.

Specify number of cores >= 1.

Load OPENMPI MPI library or IntelMPI

Please do the Exercise;  
Ex-3 in Repo

# How to submit GPU jobs?

```
#!/bin/bash
```

```
#Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --partition=gpu2
```

```
#SBATCH --gres=gpu:1
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=8
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

```
./my_executable
```

Specify partition `gpu2`

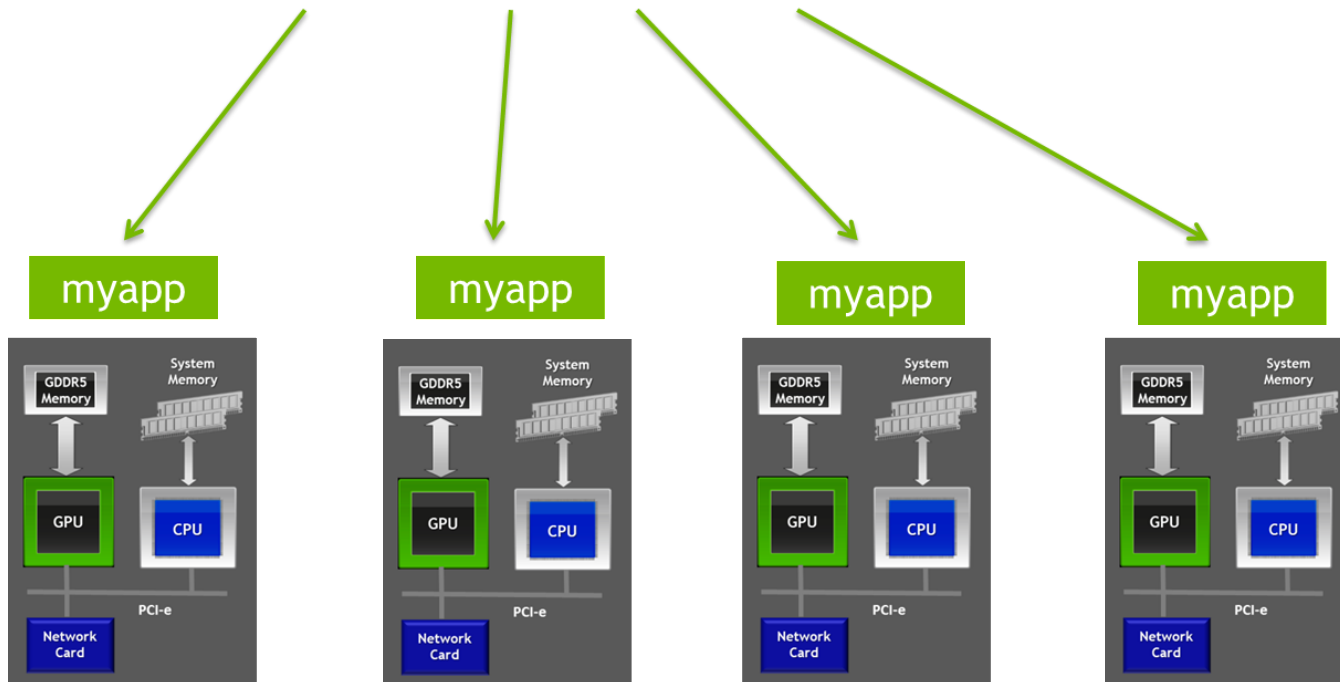
Specify number of gpus, like `gpu:1`

Please do the Exercise;  
Ex-4 in Repo

# How to submit MPI + GPU jobs?

## CUDA Aware MPI

```
mpirun -np 4 ./myapp <args>
```



# How to submit MPI + GPU jobs?

## Compilation

```
#!/bin/bash
module load openmpi/3.1.2
module load cuda/10.1
# Compiling the device code
nvcc -c dev.cu
# Compiling the host code
mpicc -c hostname.c
# Linking the host and device code
mpicc -o HostMap dev.o hostname.o -lcudart
# Submitting the job as batch script
sbatch mpijob.sh
```

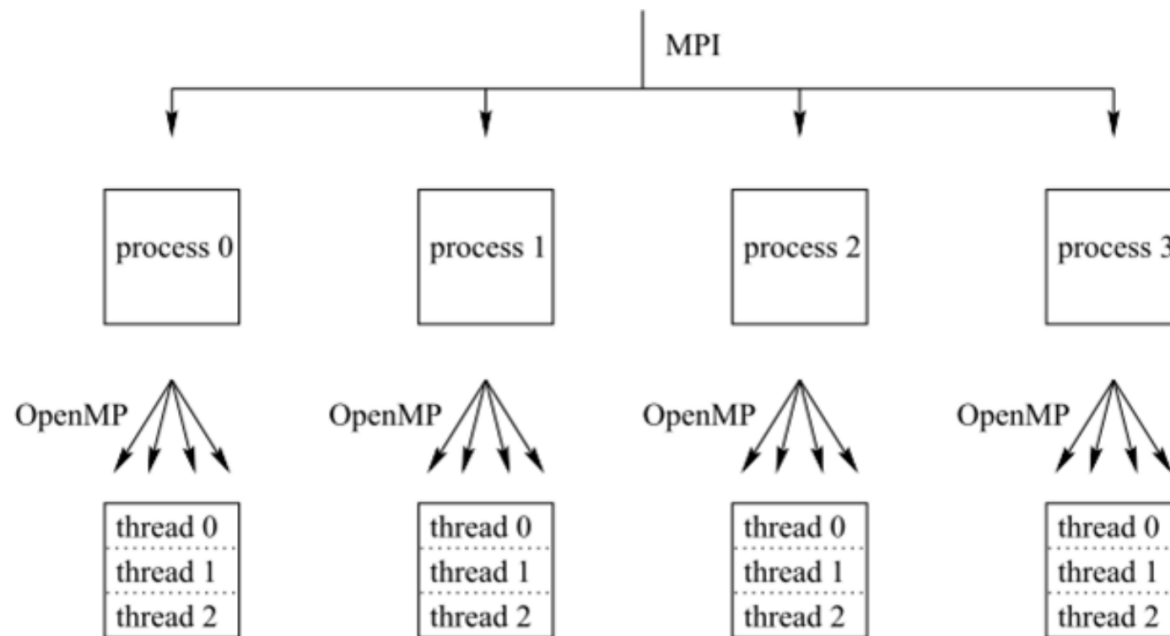
## Job Submission

```
#!/bin/bash
#SBATCH -t 00:30:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=4
#SBATCH --partition=gpu2
#SBATCH --gres=gpu:2
#SBATCH --job-name=MyJob
#SBATCH --output=MyJob-%j.out
#SBATCH --error=MyJob-%j.err
#SBATCH --qos=stafftest
mpirun ./HostMap
```



Please do the Exercise;  
Ex-5 in Repo

# How to submit MPI + OpenMPI jobs?



```
#!/bin/bash
#SBATCH --job-name=hybrid
#SBATCH --output=hybrid_%j.out
#SBATCH --error=hybrid_%j.err
#SBATCH --time=00:10:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=8
#SBATCH --partition=broadwl
#SBATCH --constraint=edr
#SBATCH --qos=stafftest

# Load the default OpenMPI module.
module load openmpi

# Set OMP_NUM_THREADS to the number of CPUs per task we asked for.
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

mpirun ./mpomp
```

Please do the Exercise;  
Ex-5a in Repo

# How to submit array based jobs

```
#!/bin/bash
```

```
# Job Name
```

```
#SBATCH --job-name=arrayjob
```

```
# Walltime requested
```

```
#SBATCH --time=0:10:00
```

```
#Add partition
```

```
#SBATCH --partition=broadwl-lc
```

```
# Provide index values (TASK IDs)
```

```
#SBATCH --array=1-16
```

```
# Use '%A' for array-job ID, '%J' for job ID and '%a' for task ID
```

```
#SBATCH --error=maths%A-%a.err
```

```
#SBATCH --output=maths%A-%a.out
```

```
# single core
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
# Use the $SLURM_ARRAY_TASK_ID variable to provide different inputs for each job
```

```
input=$((SLURM_ARRAY_TASK_ID*1000+2))
```

```
echo "Running job array number: "$SLURM_ARRAY_TASK_ID "input " $input
```

Please do the Exercise;  
Ex-6 in Repo

# How to submit dependent jobs

## SLURM Rule:

*sbatch --dependency=type:job\_id jobfile*

```
# first job - no dependencies
jobID_1=$(sbatch preprocessing.sh | cut -f 4 -d ' ')

# second job - depends on job1
jobID_2=$(sbatch --dependency=afterok:$jobID_1 analysis.sh | cut -f 4 -d ' ')

# third job - depends on job2
sbatch --dependency=afterany:$jobID_2 postprocessing.sh
```

# How to submit dependent jobs

after	This job can begin execution after the specified jobs have begun execution
afterany	This job can begin execution after the specified jobs have terminated.
aftercorr	A task of this job array can begin execution after the corresponding task ID in the specified job has completed successfully
afternotok	This job can begin execution after the specified jobs have terminated in some failed state
afterok	This job can begin execution after the specified jobs have successfully executed
singleton	This job can begin execution after any previously launched jobs sharing the same job name and user have terminated

Please do the Exercise;  
Ex-7 in Repo



# How to submit Parallel batch jobs

```
#!/bin/sh
```

```
#SBATCH --time=01:00:00
```

```
#SBATCH --partition=broadwl
```

```
#SBATCH --ntasks=28
```

```
#SBATCH --mem-per-cpu=2G # NOTE DO NOT USE THE --mem= OPTION
```

```
# Load the default version of GNU parallel. module load parallel
```

```
# set the max number of processes (which determine the max per processor)
```

```
ulimit -u 10000
```

```
# This specifies the options used to run srun. The "-N1 -n1" options are
```

```
# used to allocates a single core to each task.
```

```
srun="srun --exclusive -N1 -n1"
```

```
#Run GNU parallel
```

```
parallel="parallel --delay 0.2 -j $SLURM_NTASKS --joblog runtask.log --resume"
```

```
# Run a script, runtask.sh, using GNU parallel and srun.
```

```
$parallel "$srun ./runtask.sh arg1:{1} > runtask.sh.{1}" ::: {1..128}
```

```
# Note that if your program does not take any input, use the -n0 option to call the parallel command: # # $parallel -n0
```

```
"$srun ./run_noinput_task.sh > output.{1}" ::: {1..128}
```

# How to submit Parallel batch jobs

Please do the Exercise;  
Ex-8 in Repo

# What resources should I ask for?

This depends on the code you are running

# What resources should I ask for?

This depends on the code you are running

# What resources should I ask for?

This depends on the code you are running

## Nodes/Cores

- Question: is your code parallel? You will need to find out if your code can
  - Run on multiple cores? Run across multiple nodes?
  - Check if your code is threaded, multiprocessor, MPI
- Question: Is your code serial?
  - This means it can only make use of one core

# What resources should I ask for?

This depends on the code you are running

## Wall Time

Make an estimate of your job run and add a bit.

e.g. if think your code will take an hour, give it 1 hour and 30 min

- If your job runs out of time, your job will be killed, so
- be accurate with your estimate without going below.

# What resources should I ask for?

This depends on the code you are running

## Memory

For memory, this can take some trial and error. You can ask for a lot, then measure your usage. If you have asked for more memory and then reduce your memory with the next job.

- To ask for all the memory available on a node, use `#SBATCH --mem=0`

# What if I need an entire node or specific features?

Add this in your batch script

```
#SBATCH --exclusive
```

Add this in your batch script

```
#SBATCH --constraint=v100
```



# How do I know the features of the node to use with #SBATCH -constraint?

**nodestatus**

nodestatus

\*\*\*  
Status of nodes:  
\*\*\*

Features

Features

				STATUS	CORES IN USE	MEM IN USE	PURPOSE
midway2-0002	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	16 57.1%	24GB 41.9%	broadwl
midway2-0003	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28 100%	5GB 9%	broadwl
midway2-0004	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28 100%	15GB 26.7%	broadwl
midway2-0005	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28 100%	7GB 12.4%	broadwl
midway2-0006	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	25 89.2%	16GB 28.7%	broadwl
midway2-0007	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	16 57.1%	4GB 8.2%	broadwl
midway2-0008	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	19 67.8%	19GB 33.1%	broadwl

# What resources should I ask for?

This depends on the code you are running

**GPUs** If your code is built to use gpus you can submit to the gpu partition. To request 1 gpu:

```
#SBATCH -p gpu2 --gres=gpu:1
```

# What resources did my job actually use?

It is good practice to occasionally check what resources your job is using. For example if you are going to be submitting hundreds of similar jobs, you may save yourself a lot of waiting time in the queue by checking that you are not over requesting resources. Midway2 has a script to display the resources a job used:

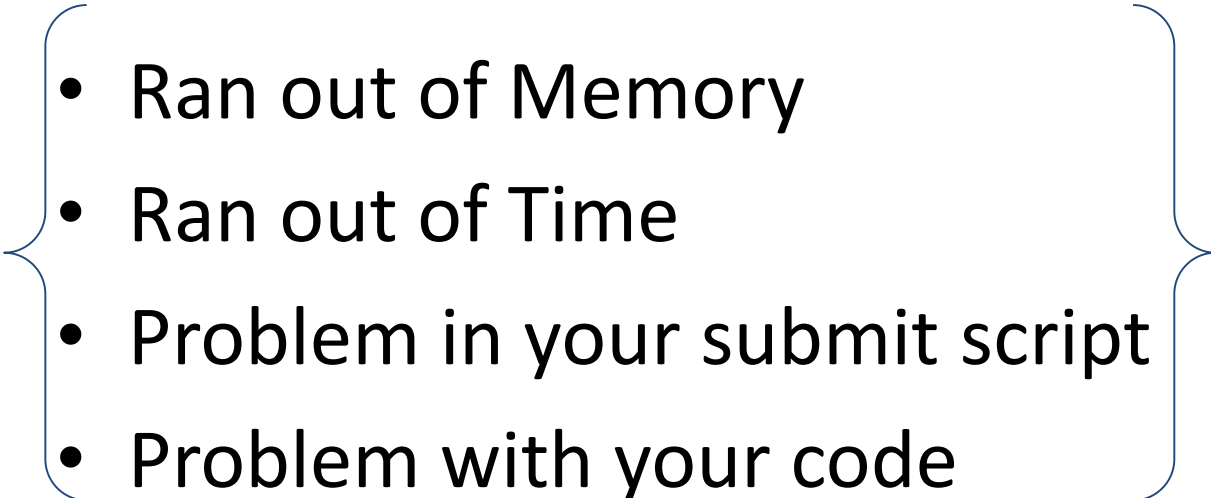
```
jobinfo -j 99999
```

**Replace 99999 with the job ID of the job you are interested in.**

# Why did my job fail?

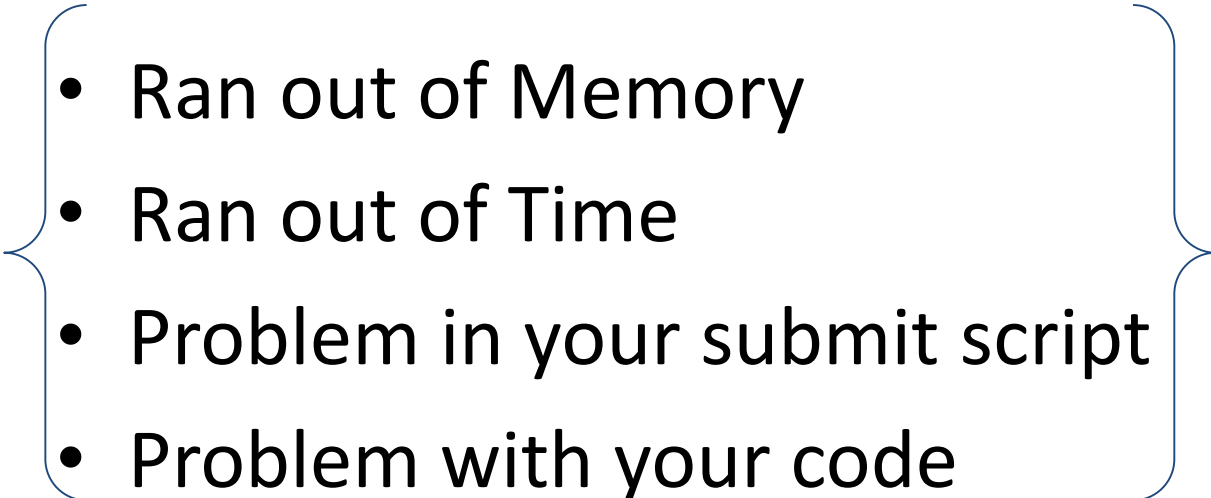
- Ran out of Memory
- Ran out of Time
- Problem in your submit script
- Problem with your code
- Node failure

# Why did my job fail?

- 
- Ran out of Memory
  - Ran out of Time
  - Problem in your submit script
  - Problem with your code
  - Node failure

You can fix these

# Why did my job fail?

- 
- Ran out of Memory
  - Ran out of Time
  - Problem in your submit script
  - Problem with your code
- Node failure

You can fix these

=> For node failure – please email  
[help@rcc.uchicago.edu](mailto:help@rcc.uchicago.edu), if this happens

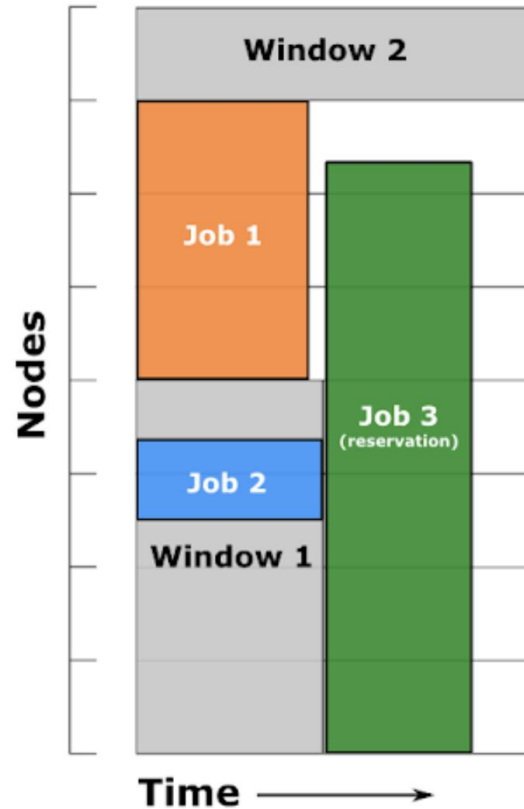
# Job Priority

- Priority is calculated using a **Fairshare** algorithm
- Fairshare is function of
  - Requested wall clock, memory, nodes/cores, etc.
  - Length of time in queue
  - Number of jobs in a time window and per PI group
  - Backfill
  - Etc.

# Job Priority

## Backfill

This blue line represents all the cores on Midway



the x axis is time

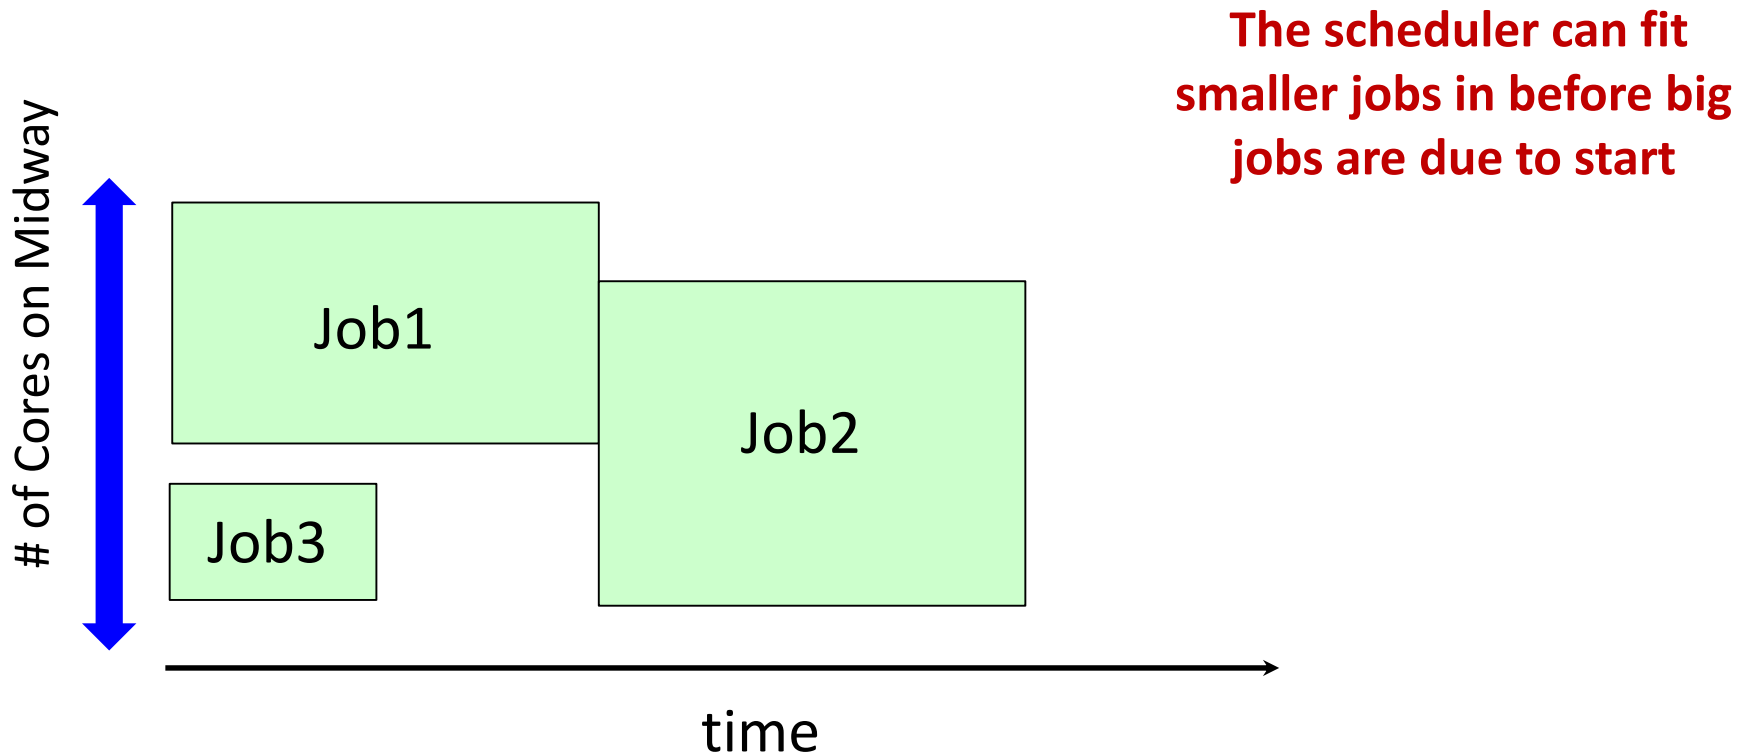
time



# Job Priority

## Backfill

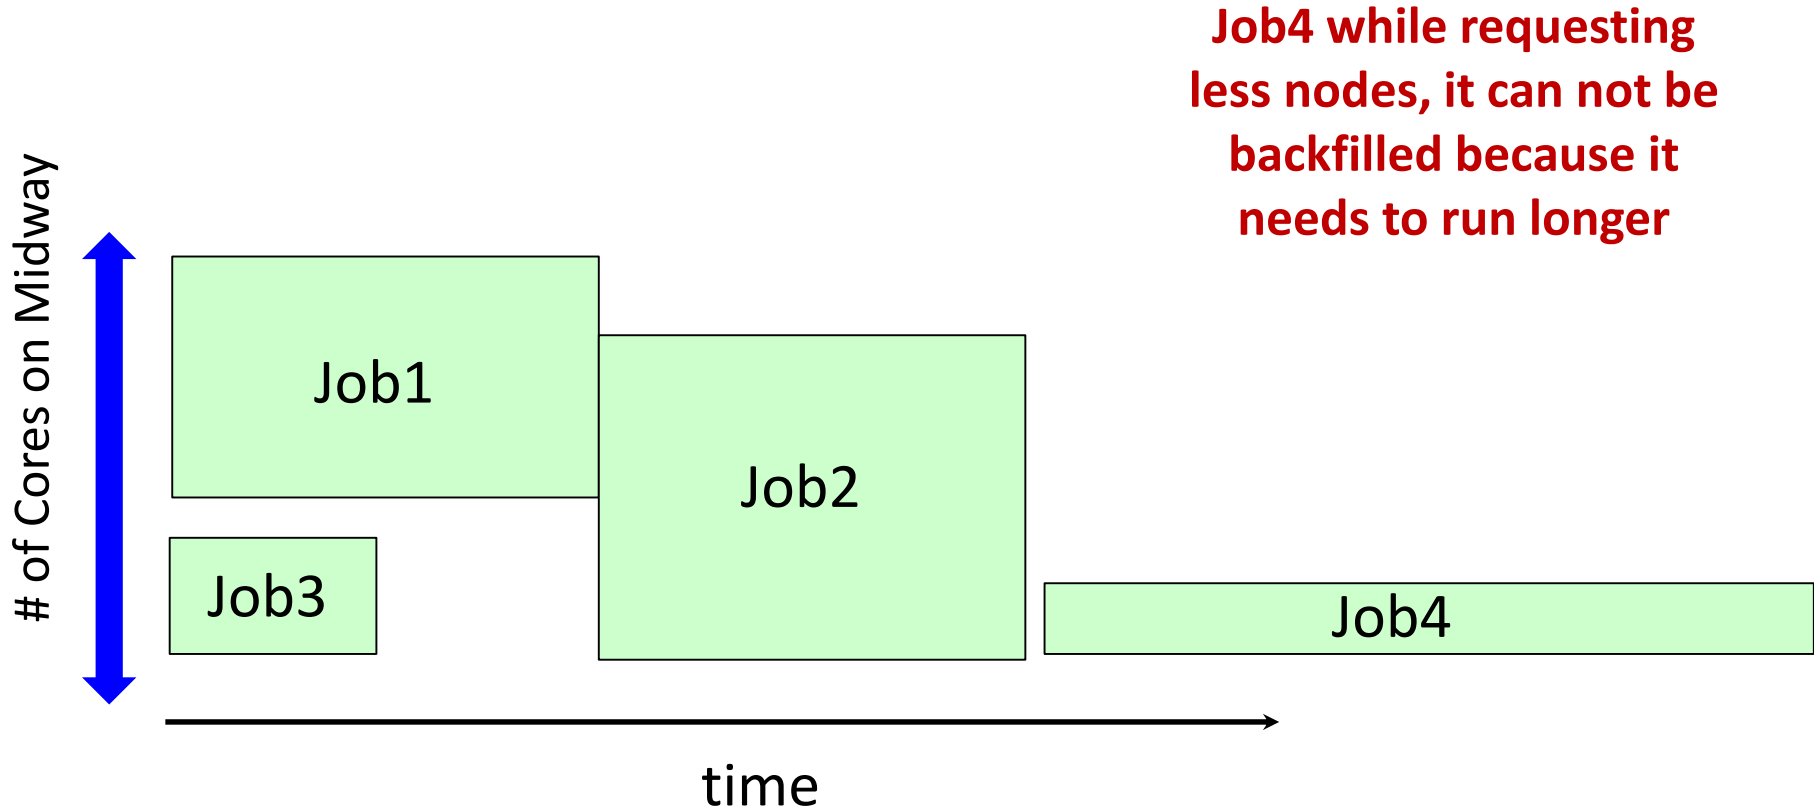
Why has someone else's job started before mine?



# Job Priority

## Backfill

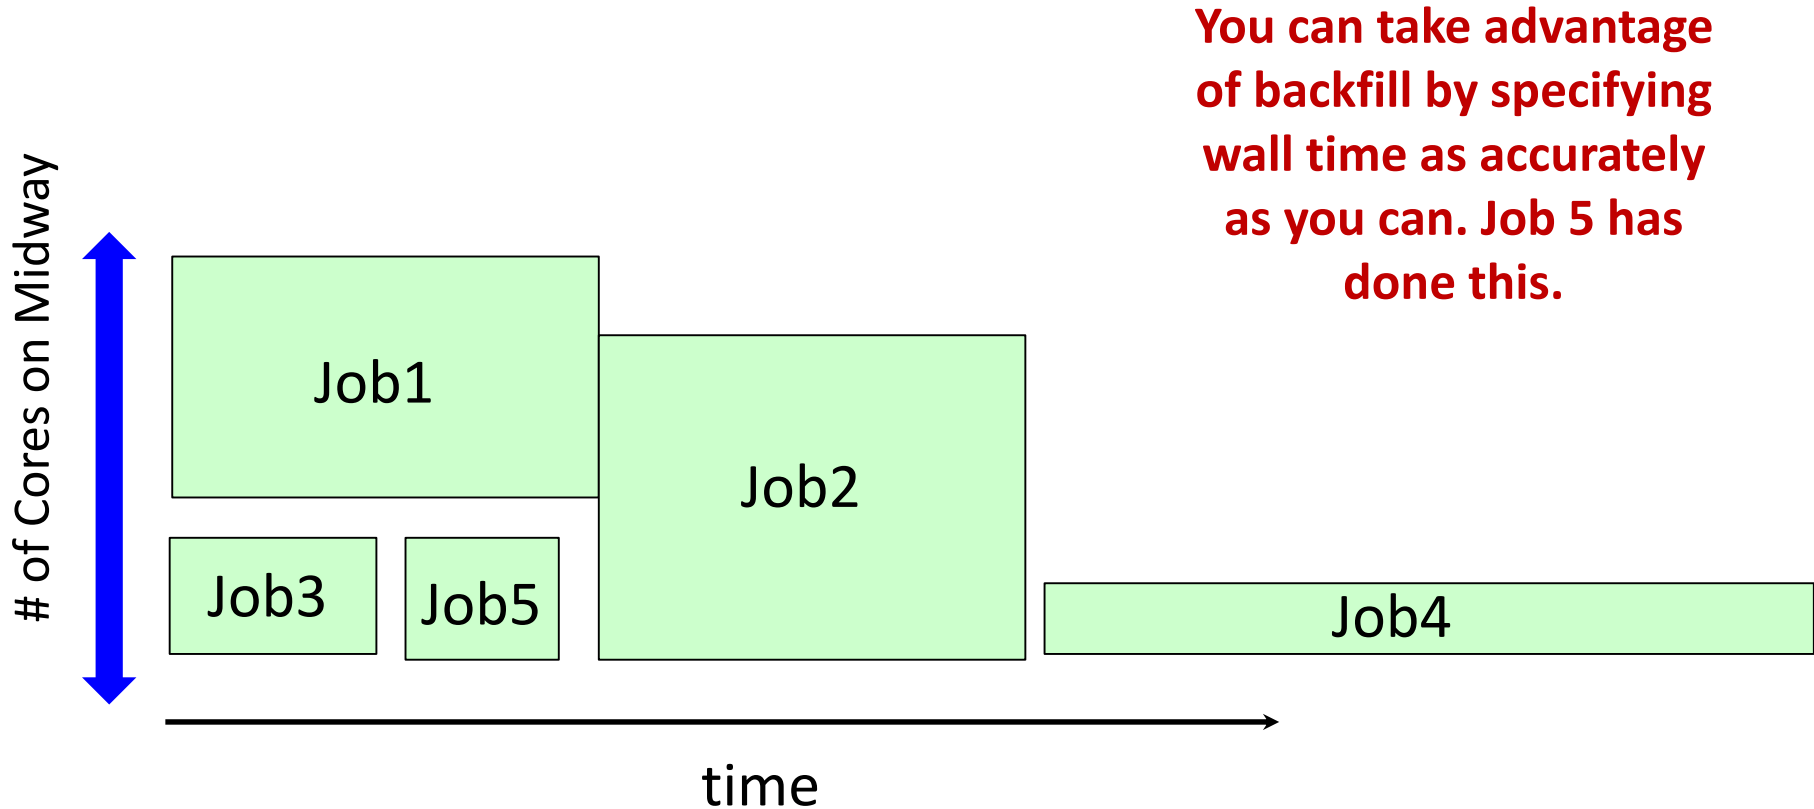
Why has someone else's job started before mine?



# Job Priority

## Backfill

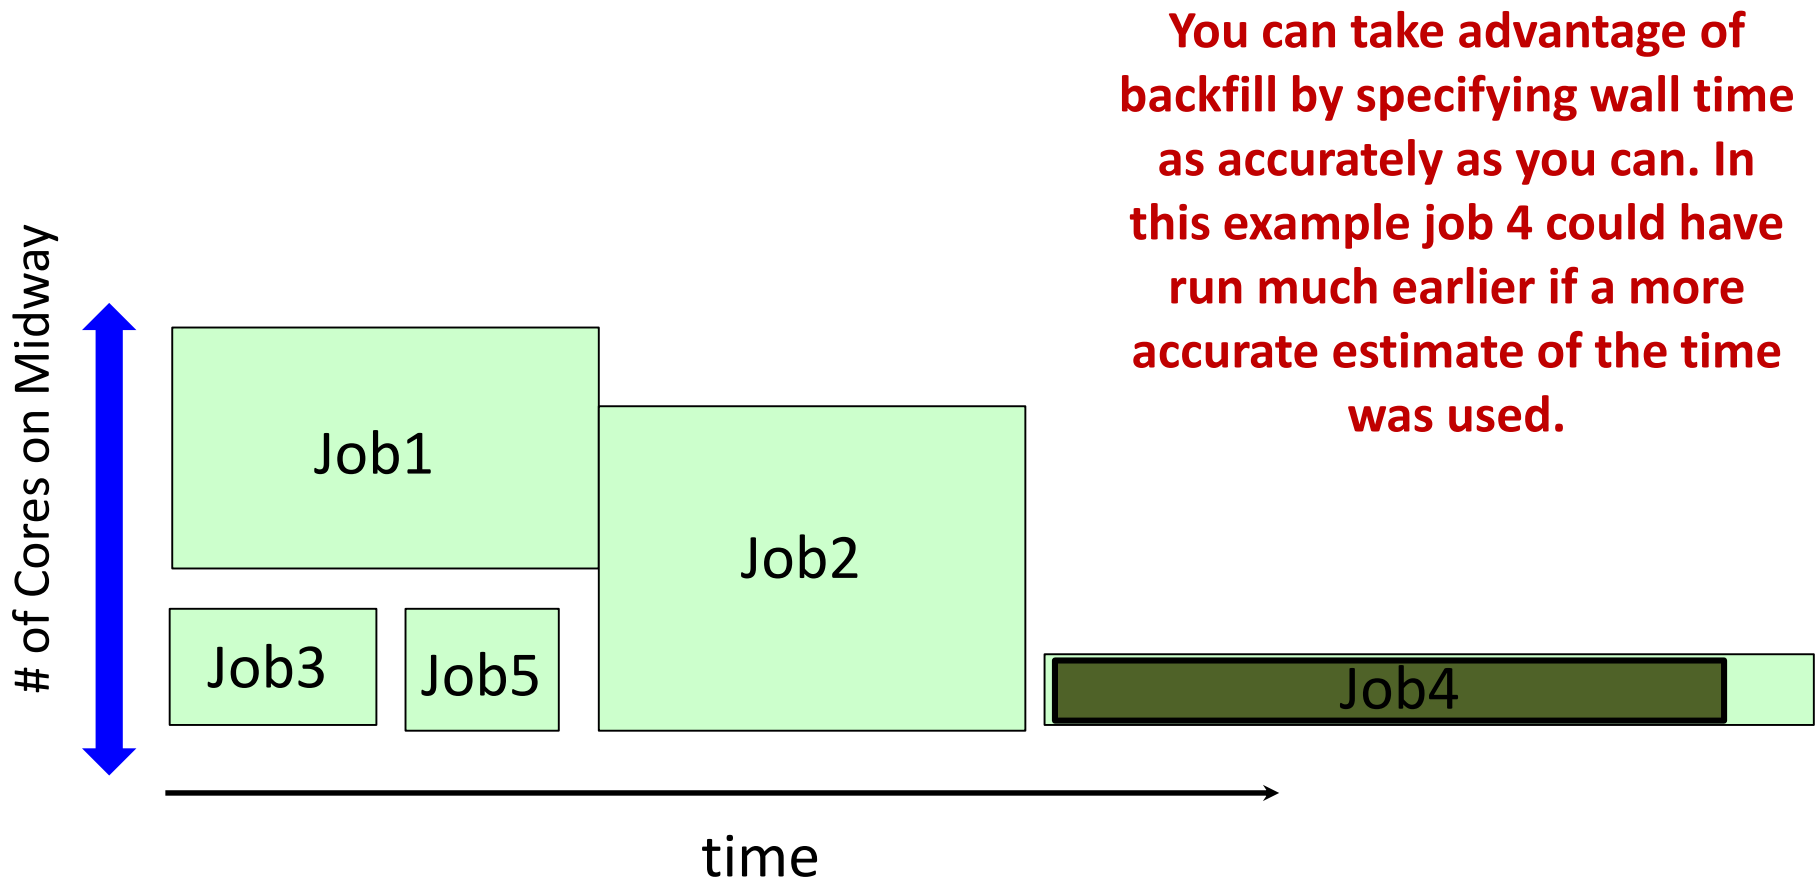
Why has someone else's job started before mine?



# Job Priority

## Backfill

Why has someone else's job started before mine?



# Job submission and monitoring

## SLURM Commands

Command	Description
<code>sbatch script.sbatch</code>	Submits <code>script.sbatch</code> job script
<code>squeue -u \$USER</code> or <code>myq</code>	Reports the status of your jobs
<code>sacct -u \$USER</code>	Displays accounting data for your job(s)
<code>scancel jobid</code>	Cancels a running job or removes it from the queue
<code>scontrol show job jobid</code> or <code>jobinfo</code>	Displays details of a running job



# Recommended online resources

- User guide on running jobs on Midway
  - <https://rcc.uchicago.edu/docs/running-jobs/index.html>
- Details Slurm documentation
  - <https://slurm.schedmd.com/sbatch.html>
- SLURM Cheat Sheet
  - <https://slurm.schedmd.com/pdfs/summary.pdf>

# RCC Help

## Contact:

A bottom–up  
approach to  
supporting  
research  
computing

- By email: [help@rcc.uchicago.edu](mailto:help@rcc.uchicago.edu)
- Web: [rcc.uchicago.edu](http://rcc.uchicago.edu)
- Phone: 773-795-2667
- In person:
  - 5607 S Drexel Avenue
  - Regenstein – Room 216
- Workshops and Tutorials:
- <http://rcc.uchicago.edu/services/training.html>

**Thank You!**